

Decidability results of Communicating Finite State Machines

G. Namratha Reddy

Chennai Mathematical Institute, India

<http://www.cmi.ac.in/~namratha>

Abstract

Models of concurrent finite state processes that communicate over queues allow different reachability results based on the underlying topology. The reachability is known to be undecidable for a general topology, but when we restrict to acyclic topologies it becomes decidable. The paper provides a proof of how it can be decided by reducing the problem to checking emptiness of a certain regular language.

2012 ACM Subject Classification Theory of computation → Distributed computing models

Keywords and phrases Distributed Systems, Reachability, Finite state machines with FIFO queues, Network of concurrent processes

1 Introduction

The concurrent nature of distributed systems makes their verification challenging. Formal methods employ mathematical models to reason about them. Communicating finite state machines (CFSM) over queues has been a standard model for modelling concurrent non-recursive programs.

Defining correctness of a concurrent program is also non-trivial. One idea is to define it via a safety property, which states that a bad state should never be reached in the program. This translates to asking whether it is possible to reach this bad state from the initial configuration in this model. The answer to this question depends on the model we consider. It turns out that for CFSMs with queues, it is undecidable to know if a state is reachable, but when restricted to certain class of acyclic topologies it becomes decidable.

2 Model

2.1 Topology

We describe our topology using the tuple

$$\mathcal{T} = (\mathcal{P}, \mathcal{C}, r, w),$$

where

\mathcal{P} is the set of processes;

\mathcal{C} is the set of channels;

$r : \mathcal{C} \rightarrow \mathcal{P}$, reader function that assigns to each channel a process that can read from it;

$w : \mathcal{C} \rightarrow \mathcal{P}$, writer function that assigns to each channel a process that can write to it.

Channels

Each channel $c \in \mathcal{C}$ has a message alphabet Γ_c . We assume channel alphabets are disjoint. If $w \in \Gamma_c^*$, we mean the first letter of w is at the head of the channel c .

Operations on the channel:

- $c!m$ is a send operation. It appends the message m to the end of channel c .
- $c?m$ is a read operation. It receives the message m from the head of channel c (enabled only if m is at the head of channel c).

42 Processes

43 Each process $p \in \mathcal{P}$ is modelled as a finite state transtion system

$$44 \quad TS_p = (Q_p, \Delta_p, s_p),$$

45 where

46 Q_p is a finite set of states;
 47 $\Delta_p \subseteq Q_p \times O_p \times Q_p$ is a finite set of transitions, where,
 48 $O_p = \{c?m \mid r(c) = p, m \in \Gamma_c\} \cup \{c!m \mid w(c) = p, m \in \Gamma_c\};$
 49 s_p is the initial state.

50 2.2 Configuration Graph

51 A configuration (α, β) of the system is a pair describing the states of all processes (control
 52 state) and the contents of all channels (channel state).

$$53 \quad \alpha \in \prod_{i \in \mathcal{P}} Q_i, \beta \in \prod_{c \in \mathcal{C}} \Gamma_c^*$$

54 For any $p \in \mathcal{P}$, $\alpha(p)$ gives the state of process p , and for any $c \in \mathcal{C}$, $\beta(c)$ gives the
 55 content of channel c . We define $\alpha[p \leftarrow q]$ to be the control state α' such that $\alpha'(p) = q$,
 56 and $\alpha'(p') = \alpha(p')$ if $p' \neq p$. Similarly, we define $\beta[c \leftarrow w] = \beta'$ such that $\beta'(c) = w$, and
 57 $\beta'(c') = \beta(c')$ if $c' \neq c$.

58 Vertices in the configuration graph $(G = (V, E))$ are configurations, and the edges are
 59 labelled with channel operations. The edges need to satisfy the following conditions:

60 Let $(\alpha, \beta), (\alpha', \beta') \in V, c \in \mathcal{C}$, then

- 61 ■ $(\alpha, \beta) \xrightarrow{c!m} (\alpha', \beta')$ if
 - 62 ■ $\alpha' = \alpha[w(c) \leftarrow \alpha'(w(c))]$
 - 63 ■ $(\alpha(w(c)), c!m, \alpha'(w(c))) \in \Delta_{w(c)}$
 - 64 ■ $\beta' = \beta[c \leftarrow \beta(c).m]$
- 65 ■ $(\alpha, \beta) \xrightarrow{c?m} (\alpha', \beta')$ if
 - 66 ■ $\alpha' = \alpha[r(c) \leftarrow \alpha'(r(c))]$
 - 67 ■ $(\alpha(r(c)), c?m, \alpha'(r(c))) \in \Delta_{r(c)}$
 - 68 ■ $\beta = \beta'[c \leftarrow m.\beta'(c)]$

69 A run ρ is a path in this graph. $c_1 \xrightarrow[G]{*} c_2$ represents a run starting in c_1 and ending in
 70 c_2 . $\rho \downarrow_p$ is the projection of ρ to process p , i.e. only transitions of process p are considered.
 71 Similarly, we can define projection to a set of processes.

72 3 The Reachability Problem

73 A target control state α_t is said to be reachable if there is a run $c_0 \xrightarrow{*} c_t$, where,
 74 $c_0 = (\alpha_0, \beta_0)$, $\alpha_0(p) = s_p \forall p \in \mathcal{P}$, and $\beta_0(c) = \epsilon \forall c \in \mathcal{C}$; $c_t = (\alpha_t, \beta)$ for some β

75 The reachability problem asks whether a given target control state is reachable. We know
 76 that the problem is undecidable in general, because if there is a loop in the topology then we
 77 can simulate a queue machine and state reachability is undecidable for queue machines.

78 So we ask the question of whether it is decidable for acyclic topologies. What we mean
 79 by acyclic topologies is consider the network topology and ignore the direction of the edges,
 80 we get an undirected graph and this graph should have no cycles.

81 Tree topology

82 Given an acyclic topology the reachability problem can be reduced to one in which the
 83 target state is reached only if the queue is empty [1]. Using this reduction our reachability
 84 problem reduces to reaching a channel state β_t with every channel being empty i.e. we take
 85 every channel content to be ϵ in the target state.

86 We further reduce our topology to another isomorphic topology that looks like a tree,
 87 where every process has one incoming edge (except the root) i.e every process can read from
 88 one channel but write to multiple channels [1]. Henceforth we work with a tree topology.
 89 Although the root node doesn't read from any channel, we add a dummy channel which
 90 is initialized to ϵ for consistency with other nodes. This way every node reads from some
 91 channel and we define a function $\iota(p)$, which gives the channel that process p reads from.
 92 We define $G_p = (V_p, E_p)$ to be the induced subgraph of the configuration graph G where V_p
 93 are the nodes reachable from p in the tree.

94 Languages

95 Given the target control state α_t to be reached, we define for all $p \in \mathcal{P}$ the language
 96 $L_p = L(A_p)$. We get the finite automation A_p by taking the transition system TS_p and
 97 replacing the operation in the transition with just the message and having final state as $\alpha_t(p)$.
 98 $A_p = (Q_p, \Sigma_p, \delta_p, s_p, \alpha_t(p))$, where, $\Sigma_p = \bigcup_{c \in \mathcal{C}} \{\Gamma_c : r(c) = p \text{ or } w(c) = p\}$, $\delta_p = \{(q, m, q') :$
 99 $(q, c?m, q') \in \Delta_p \text{ or } (q, c!m, q') \in \Delta_p\}$. $\pi_c(w)$ is the word we get by deleting all letters not in
 100 Γ_c . This can be extended to languages as well i.e. $\pi_c(L) = \{\pi_c(w) : w \in L\}$.

101 We define for all $p \in \mathcal{P}$, L_p^e as follows :

- 102 ■ if p is a leaf node, $L_p^e = L_p \cap \Gamma_{\iota(p)}^*$,
- 103 ■ if p is a non-leaf node, and children of p is given by the set $k(p) = \{k_1, k_2, \dots, k_l\}$,
 104 $L_p^e = L_p \cap \text{shuffle}(\{\pi_{\iota(i)}(L_i^e) : i \in k(p)\} \cup \{\Gamma_{\iota(p)}^*\})$

105 If there is run ρ in G then by projecting it to p we get $\rho \downarrow_p$, we can take this run and
 106 erase the channel name to get a corresponding run in A_p . Conversely, if we have a run in A_p
 107 labelled with w and start in a configuration where channel $\iota(p)$ has the word $\pi_{\iota(p)}(w)$ then
 108 we can get the corresponding run in G .

109 ► **Lemma 1.** *Let p be the root of the tree, then*

$$110 \quad w \in L_p^e \implies (\alpha_0, \beta_0[\iota(p) \leftarrow \pi_{\iota(p)}(w)]) \xrightarrow[G_p]{*} (\alpha'_t, \beta_t),$$

111 where $\alpha'_t(p) = \alpha_t(p)$ for $p \in V_p$ and $\alpha'_t(p) = \alpha_0(p)$ for $p \notin V_p$

112 **Proof.** By induction on the no. of nodes in the tree graph G_p .

113 Base case: No. of nodes is 1. Let this node be p . $w \in L_p^e \implies w \in L_p$, so there is a run
 114 $s_p \xrightarrow[A_p]{*} \alpha_t(p)$. Since p is a leaf node it doesn't write to any channel, $\pi_{\iota(p)}(w) = w$. Therefore,
 115 in G_p we get the run $(\langle s_p \rangle, \langle \pi_{\iota(p)}(w) \rangle) \xrightarrow[G_p]{*} (\langle \alpha_t(p) \rangle, \langle \epsilon \rangle)$

116 Induction: Let p be a non-leaf node and children of p be $k(p) = \{k_1, k_2, \dots, k_l\}$. We have
 117 $w \in L_p \cap \text{shuffle}(\{\pi_{\iota(i)}(L_i^e) : i \in k(p)\} \cup \{\Gamma_{\iota(p)}^*\}) \implies w \in \text{shuffle}(w_1, \dots, w_l, \pi_{\iota(p)}(w))$ where
 118 $w_i \in \pi_{\iota(i)}(L_i^e)$. By induction hypothesis we have a run $\rho_i : (\alpha_0, \beta_0[\iota(i) \leftarrow w_i]) \xrightarrow[G_i]{*} (\alpha'_t, \beta_t)$
 119 for each $i \in k(p)$. Since $w \in L_p$ we have the run $\rho : s_p \xrightarrow[A_p]{*} \alpha_t(p)$. Using ρ and the runs for
 120 each i , ρ_i , we can construct a run in G_p as follows: $(\alpha_0, \beta_0[p \leftarrow \pi_{\iota(p)}(w)]) \xrightarrow[p \text{ moves}]{*} (\alpha_0[p \leftarrow$
 121 $\alpha_t(p)], \beta_0[\{i \leftarrow w_i : i \in k(p)\}]) \xrightarrow[k_1 \text{ moves}]{*} \dots \xrightarrow[k_l \text{ moves}]{*} (\alpha'_t, \beta_t)$ ◀

122 ► **Lemma 2.** *Let p be the root of the tree, then*

$$123 \quad (\alpha_0, \beta_0[\iota(p) \leftarrow \omega]) \xrightarrow{G_p^*} (\alpha'_t, \beta_t) \implies \exists w \in L_p^e \text{ such that } \pi_{\iota(p)}(w) = \omega,$$

124 *where $\alpha'_t(p) = \alpha_t(p)$ for $p \in V_p$ and $\alpha'_t(p) = \alpha_0(p)$ for $p \notin V_p$*

125 **Proof.** Induction on the no. of nodes in the tree.

126 Base case: No. of nodes is 1. Let this node be p . We have $(\langle s_p \rangle, \langle \omega \rangle) \xrightarrow{G_p^*} (\langle \alpha_t(p) \rangle, \langle \epsilon \rangle)$

127 implies we have the run $s_p \xrightarrow{A_p^\omega} \alpha_t(p) \implies \omega \in L_p^e$ and we can take w as ω itself since

$$128 \quad \pi_{\iota(p)}(\omega) = \omega.$$

129 Induction: Let p be a non-leaf node and children of p is given by the set $k(p) =$
 130 $\{k_1, k_2, \dots, k_l\}$. We have the run $\rho : (\alpha_0, \beta_0[\iota(p) \leftarrow \omega]) \xrightarrow{G_p^*} (\alpha'_t, \beta_t)$. We can construct another

131 run ρ' in which all the p transitions are taken first and no other process moves until p finishes.
 132 Once p finishes let channel state be β' where $\beta'(\iota(k_i)) = \omega_i$ and the rest of the channels have

133 ϵ .

134 Take the corresponding run in A_p and let that be $s_p \xrightarrow{A_p^{w'}} \alpha_t(p)$

135 We can extract the following run from ρ' by taking only transitions of V_{k_i}

$$136 \quad (\alpha, \beta_0[\iota(k_i) \leftarrow \omega_i]) \xrightarrow{\text{only nodes in } V_{k_i} \text{ move}} (\alpha'_t, \beta_t). \text{ By induction hypothesis we have some}$$

137 word $w_i \in L_{k_i}^e$ and $\omega_i = \pi_{\iota(k_i)}(w_i) \implies \omega_i \in \pi_{\iota(k_i)} L_{\iota(k_i)}^e$

138 We have $w' \in \text{shuf fle}(\omega_1, \dots, \omega_l, \omega) \implies w' \in L_p^e$ and $\pi_{\iota(p)}(w') = \omega$

139 ◀

140 ► **Theorem 3.** *The Reachability problem is decidable for CFSMs with queues over acyclic*
 141 *topologies.*

142 **Proof.** We use the reductions stated above to get a tree topology. Let p be the root of the
 143 tree. From Lemma 1 we know that if $w \in L_p^e$ then there is a run $(\alpha_0, \beta_0[\iota(p) \leftarrow \pi_{\iota(p)}(w)]) \xrightarrow{G_p^*}$
 144 (α'_t, β_t) . Since the root node doesn't read from a channel we can keep it empty and take the
 145 same transitions to get $\rho : (\alpha_0, \beta_0) \xrightarrow{G_p^*} (\alpha_t, \beta_t)$. Conversely, if we have such a run ρ , using
 146 Lemma 2 and substituting ϵ for ω we get a $w \in L_p^e$. Therefore, we can reduce reachability to
 147 emptiness of the language L_p^e which is regular and hence decidable.

148 ◀

149 — References —

- 150 1 La Torre, S., Madhusudan, P., Parlato, G. (2008). Context-Bounded Analysis of Concurrent
 151 Queue Systems. In: Ramakrishnan, C.R., Rehof, J. (eds) Tools and Algorithms for the
 152 Construction and Analysis of Systems. TACAS 2008. Lecture Notes in Computer Science, vol
 153 4963. Springer, Berlin, Heidelberg. https://doi.org/10.1007/978-3-540-78800-3_21